

[Return to Classroom](#)[DISCUSS ON STUDENT HUB](#)

Operationalizing Machine Learning

REVIEW

CODE REVIEW

HISTORY

Meets Specifications

Awesome 🙌🙌

Congratulations on passing this project. 🏆🏆

This project was not easy but you have done it gracefully. This is all because of your hard work and continuous evaluation. But still don't relax, keep exploring and learning from the other references. These projects are just a starting phase of learning these things but there are many more complex topics and concepts out there. So start exploring and keep learning. 💪💪

Useful Resources

- [Benchmark the Endpoint using Apache Benchmark](#)
- [Apache Benchmark Documentation](#)
- [Swagger Documentation Lesson](#)
- [Swagger REST API Documentation homepage](#)

You can also share your project on LinkedIn and ask the audience for a necessary feedback or open the project for anyone to collaborate. This way you will find many interesting connections and engagement with others.

I wish you good luck. Looking forward to your success.

For any queries, you can ask on [Knowledge Portal](#) as well.

Stay  ! Stay Safe

DON'T FORGET TO RATE MY WORK AS PROJECT REVIEWER! YOUR FEEDBACK IS VERY HELPFUL AND APPRECIATED.

Machine Learning Ops Principles

A README file is included in the project root and has:

- An overview of the project
- An Architectural Diagram
- A short description of how to improve the project in the future
- Screenshots required with a short description to demonstrate key steps
- A link to the screencast video on YouTube (or a similar alternative streaming service)
- In case you are unable to provide an audio file, you can include a written description of your script instead of audio, if you prefer. Please include it in your README file.



An overview of the project

Further Improvement

Overview

- This project is a part Machine Learning Engineer with Microsoft Azure Nanodegree program.
- In this project we implemented end-to-end machine learning project cycle in Azure ML studio. We used Azure ML to create and train the model with AutoML functionality and deploy it.
- We also created a pipeline to create the model, deploy the model and consumed the model.
- Also the end-point URL of the deployed model is exposed to end-user which they can consume it with sample data.

You have discussed the actual problem that you are solving through this. (Hint: Bank Marketing)

Architectural Diagram



An Architectural Diagram



A short description of how to improve the project in the future

[LINK to the Demo.](#)

Future Improvement

- We can try to increase the dataset to improve the model.
- We can export this model to ONNX for faster inference.
- We can include deep learning model in order get good model accuracy.

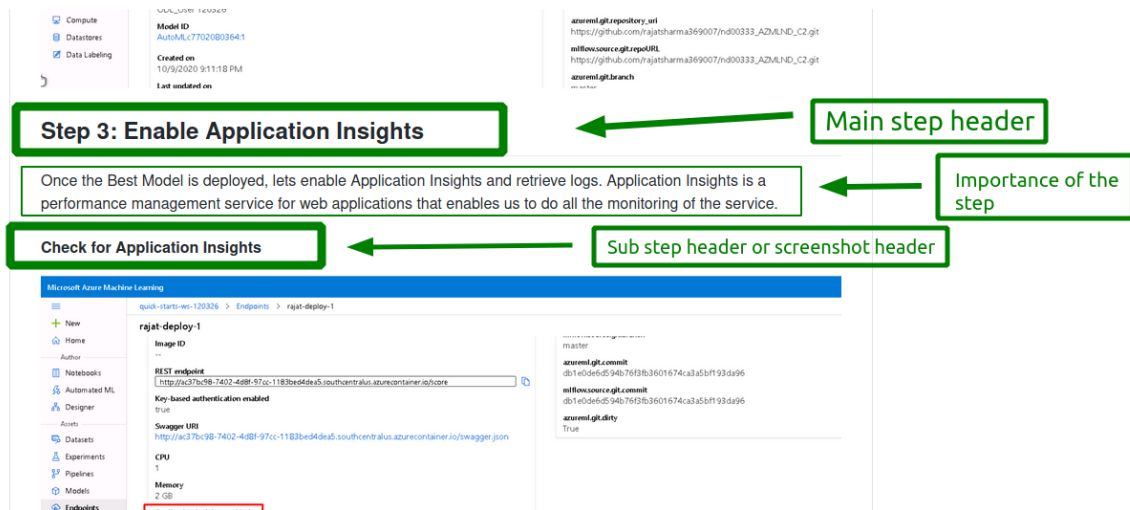
The data is also imbalanced so you can try manually balancing the data using algorithmic programs.



Screenshots required with a short description to demonstrate key steps

You have provided the screenshots but try to improve the structure.

You can make this section even better by providing h3 headers before each screenshot stating what each screenshot is about. You can also highlight the parts. For example:



✓ A link to the screencast video on YouTube (or a similar alternative streaming service)

Useful References:

- [Resource for understanding project overview link 1](#)
- [Resource for understanding project overview link 2](#)
- [Azure Automated Machine Learning Concepts](#)

The screencast should meet the following criteria:

- Screencast is 1-5 minutes in length
- Audio is clear and understandable
- Video is 1080P or higher with 16:9 aspect ratio
- text is readable

The screencast shows the entire process of the working ML application, including a demonstration of:

- Working deployed ML model endpoint.
- Deployed Pipeline
- Available AutoML Model
- Successful API requests to the endpoint with a JSON payload

Screencast is very audible, readable, and informative. Good work!

✓ Screencast is 1-5 minutes in length

- ✓
- ✓ Audio is clear and understandable
- ✓ Video is 1080P or higher with 16:9 aspect ratio
- ✓ text is readable

- ✓ Working deployed ML model endpoint.

Available at 3:44

- ✓ Deployed Pipeline

Available at 3:01

- ✓ Available AutoML Model

Available at 0:43

- ✓ Successful API requests to the endpoint with a JSON payload


Available at 3:44

Deploy model in Azure ML Studio

The submission includes screenshots of:

- "Registered Datasets" in ML Studio shows "Bankmarketing" dataset available
- The experiment is shown as completed.

- ✓ "Registered Datasets" in ML Studio shows "Bankmarketing" dataset available



Home > Datasets

Datasets

Registered datasets Dataset monitors (preview)

+ Create dataset ▾ Refresh ☆ Unregister

Name	Version	Data source	Created on	Modified on
BankMarketing	1	workspaceblobstore	Apr 13, 2021 12:24 PM	Apr 13, 2021 12:24 PM

- ✓ The experiment is shown as completed.

Run 3 ✔ Completed

[Refresh](#) [Cancel](#)

Details Data guardrails Models Outputs + logs Child runs Snapshot

Properties

- Status: ✔ Completed
- Created: Apr 14, 2021 4:00 AM
- Started: Apr 14, 2021 4:01 AM
- Duration: 21m 31.50s
- Compute target: cluster-...

Best model summary

- Algorithm name: [VotingEnsemble](#)
- Accuracy: 0.91866 [View all](#)
- Sampling: 100.00 % [?](#)
- Registered models: No registration yet
- Deploy status: No deployment yet

The submission includes screenshots of:

- Endpoints section in Azure ML Studio, showing that “Application Insights enabled” says “true”.
- Logging is enabled by running the provided `logs.py` script
- Swagger runs on localhost showing the HTTP API methods and responses for the model
- `endpoint.py` script runs against the API producing JSON output from the model.
- Apache Benchmark (ab) runs against the HTTP API using authentication keys to retrieve performance results. (optional)

✔ Endpoints section in Azure ML Studio, showing that “Application Insights enabled” says “true”.

Key-based authentication enabled
true

Swagger URI
<http://3b8c8dd1-a3da-46e0-941e-e2a11c81d804.eastus2.azurecontainer.io/swagger.json>

CPU
1.8

Memory
4 GB

Application Insights enabled
true

✔ Logging is enabled by running the provided `logs.py` script

✓ Logging is enabled by running the provided `logger.py` script

```
/usr/sbin/nginx: /azureml-envs/azureml_8e5a5a51349877e7d47c6a2872e0ebfd/lib/libcrypto.so.1.0.0: no version information available (required by /usr/sbin/nginx)
/usr/sbin/nginx: /azureml-envs/azureml_8e5a5a51349877e7d47c6a2872e0ebfd/lib/libcrypto.so.1.0.0: no version information available (required by /usr/sbin/nginx)
/usr/sbin/nginx: /azureml-envs/azureml_8e5a5a51349877e7d47c6a2872e0ebfd/lib/libssl.so.1.0.0: no version information available (required by /usr/sbin/nginx)
/usr/sbin/nginx: /azureml-envs/azureml_8e5a5a51349877e7d47c6a2872e0ebfd/lib/libssl.so.1.0.0: no version information available (required by /usr/sbin/nginx)
/usr/sbin/nginx: /azureml-envs/azureml_8e5a5a51349877e7d47c6a2872e0ebfd/lib/libssl.so.1.0.0: no version information available (required by /usr/sbin/nginx)
EdgeHubConnectionString and IOTEDGE_IOTHUBHOSTNAME are not set. Exiting...
2021-04-13T23:04:12,947946400+00:00 - iot-server/finish 1 0
2021-04-13T23:04:12,952494100+00:00 - Exit code 1 is normal. Not restarting iot-server.
Starting gunicorn 19.9.0
Listening at: http://127.0.0.1:31311 (72)
Using worker: sync
worker timeout is set to 300
Booting worker with pid: 101
SPARK_HOME not set. Skipping PySpark Initialization.
Generating new fontManager, this may take some time...
Initializing logger
2021-04-13 23:04:14,447 | root | INFO | Starting up app insights client
2021-04-13 23:04:14,448 | root | INFO | Starting up request id generator
2021-04-13 23:04:14,448 | root | INFO | Starting up app insight hooks
2021-04-13 23:04:14,448 | root | INFO | Invoking user's init function
2021-04-13 23:04:17,818 | azureml.core | WARNING | Failure while loading azureml_run_type_providers. Failed to load endpoint autorml = azureml.train.autorml.run:AutoMLRun._from_run_dto with exception cannot import name 'RunType'.
Failure while loading azureml_run_type_providers. Failed to load endpoint autorml = azureml.train.autorml.run:AutoMLRun._from_run_dto with exception cannot import name 'RunType'.
2021-04-13 23:04:18,231 | root | INFO | User's init has completed successfully
2021-04-13 23:04:18,236 | root | INFO | Skipping middleware: dbg_model_info as it's not enabled.
2021-04-13 23:04:18,236 | root | INFO | Skipping middleware: dbg_resource_usage as it's not enabled.
```

No traceback

✓ Swagger runs on localhost showing the HTTP API methods and responses for the model

Swagger
Powered by SMARTBEAR

http://localhost:8000/swagger.json

bankmarketing-prediction-deploy 1.0

http://localhost:8000/swagger.json

API specification for the Azure Machine Learning service bankmarketing-prediction-deploy

Schemes
HTTPS

default

GET /

✓ `endpoint.py` script runs against the API producing JSON output from the model.

1. We will modify the endpoint url and key in `endpoint.py` and run it to check the output.

```
(base) → starter_files git:(master) × python endpoint.py
{"result": ["no", "no"]}
(base) → starter_files git:(master) ×
```


✓

✓ Apache Benchmark (ab) runs against the HTTP API using authentication keys to retrieve performance results. (optional)

```
HTML transferred: 320 bytes
Requests per second: 1.28 [# /sec] (mean)
Time per request: 783.314 [ms] (mean)
Time per request: 783.314 [ms] (mean, across all concurrent requests)
Transfer rate: 0.32 [Kbytes/sec] received
               1.32 kb/s sent
               1.64 kb/s total

Connection Times (ms)
      min  mean[+/-sd] median  max
Connect:  201   300 253.1   207  1013
Processing: 466   483  15.3   487   516
Waiting:    465   482  14.3   484   511
Total:      668   783 264.6   695  1529

Percentage of the requests served within a certain time (ms)
 50%    695
 66%    696
 75%    731
 80%    793
 90%   1529
 95%   1529
```



Excellent work! You have also completed the optional part of benchmarking the deployed endpoint.

Useful Resources

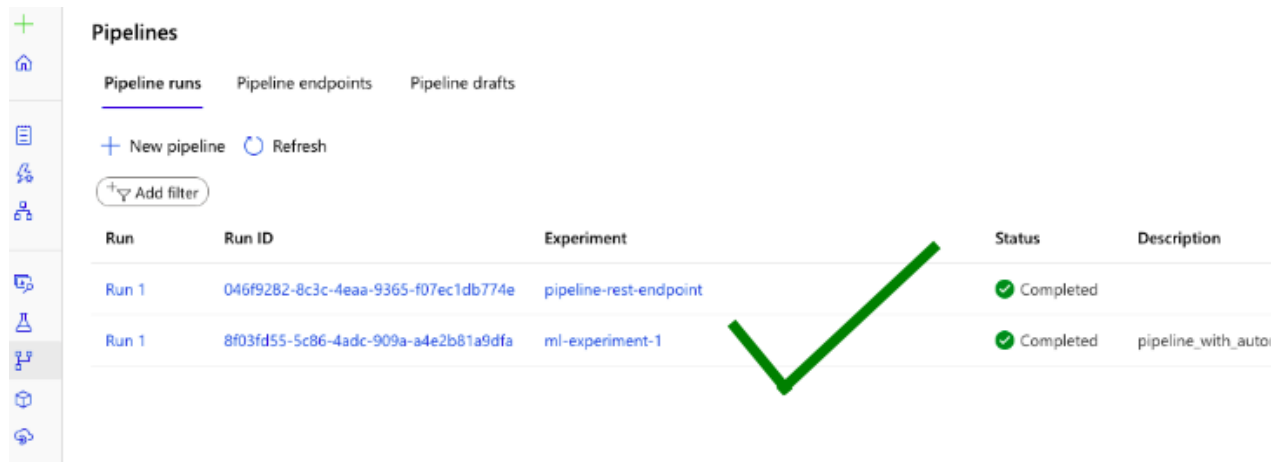
- [Benchmark the Endpoint using Apache Benchmark](#)
- [Apache Benchmark Documentation](#)
- [Swagger Documentation Lesson](#)
- [Swagger REST API Documentation homepage](#)

Publish an ML Pipeline

The submission includes screenshots of:

- The pipeline section of Azure ML studio, showing that the pipeline has been created
- The Bankmarketing dataset with the AutoML module
- The “Published Pipeline overview”, showing a REST endpoint and a status of ACTIVE

✓ The pipeline section of Azure ML studio, showing that the pipeline has been created



Pipelines

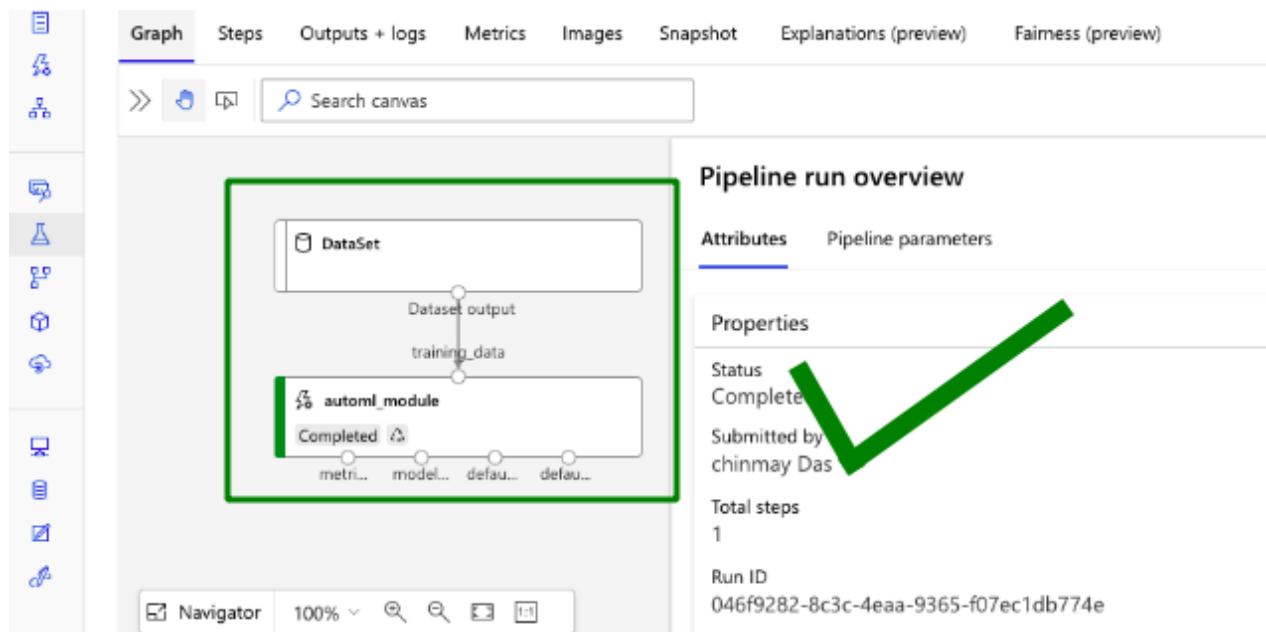
Pipeline runs Pipeline endpoints Pipeline drafts

+ New pipeline Refresh

+ Add filter

Run	Run ID	Experiment	Status	Description
Run 1	046f9282-8c3c-4eaa-9365-f07ec1db774e	pipeline-rest-endpoint	Completed	
Run 1	8f03fd55-5c86-4adc-909a-a4e2b81a9dfa	ml-experiment-1	Completed	pipeline_with_auto

✓ The Bankmarketing dataset with the AutoML module



Graph Steps Outputs + logs Metrics Images Snapshot Explanations (preview) Fairness (preview)

Search canvas

Pipeline run overview

Attributes Pipeline parameters

Properties

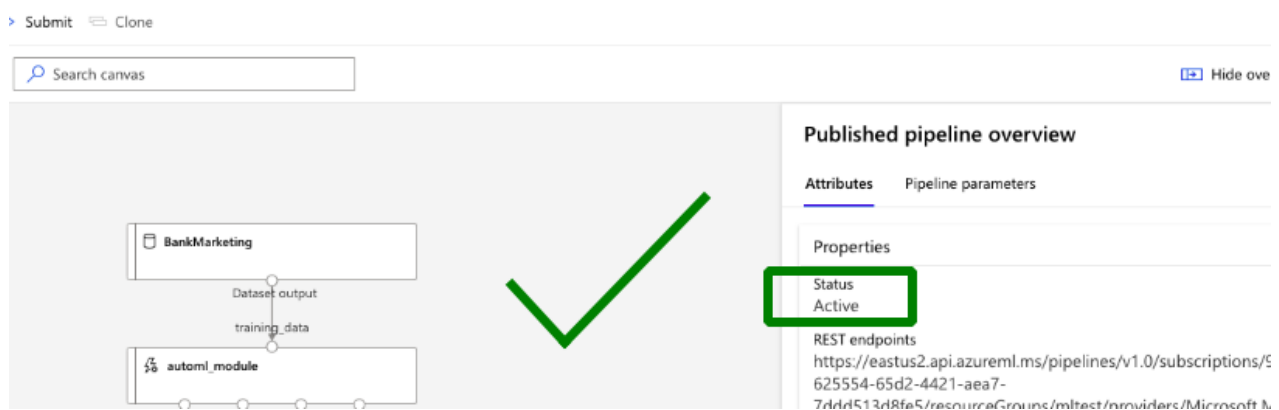
Status: Completed

Submitted by: chinmay Das

Total steps: 1

Run ID: 046f9282-8c3c-4eaa-9365-f07ec1db774e

✓ The "Published Pipeline overview", showing a REST endpoint and a status of ACTIVE



> Submit Clone

Search canvas

Published pipeline overview

Attributes Pipeline parameters

Properties

Status: Active

REST endpoints: <https://eastus2.api.azureml.ms/pipelines/v1.0/subscriptions/625554-65d2-4421-aea7-7dd4513d8fe5/resourceGroups/mltest/providers/Microsoft.ML/pipelines/046f9282-8c3c-4eaa-9365-f07ec1db774e>

metri... model... defau... defau...

rade... chineLearningServices/workspaces/mltest/PipelineRuns/Pipel...
 eSubmit/ae0c5fe3-88e6-4aff-a13f-6b1703d076ea
 Updated by
 chinmay Das

- A screenshot of the Jupyter Notebook is included in the submission showing the “Use RunDetails Widget” with the step runs

✓ A screenshot of the Jupyter Notebook is included in the submission showing the “Use RunDetails Widget” with the step runs

The screenshot shows a Jupyter Notebook interface. At the top, there's a code cell with the following code:

```
from azureml.widgets import RunDetails
RunDetails(pipeline_run).show()
```

Below the code cell, there's a 'Properties' table and an 'Output Logs' section.

Property	Value
Status	Completed
Run Start Time	13/04/2021 15:55:10
Run Duration	0:41:29
Run Id	8f03fd55-5c86-4adc-909a-a4e2b81a9dfa
Run Time	13/04/2021 16:36:39

The 'Output Logs' section shows the following logs:

```
[2021-04-13 10:25:20Z] Submitting 1 runs, first five are: bdb561c9:6070fd8a-787f-4d12-c3078c73c750
[2021-04-13 11:06:39Z] Completing processing run id 6070fd8a-787f-4d12-afce-c3078c
Run is completed.
```

A large green checkmark is overlaid on the logs.

The submission includes screenshots of:

- ML studio showing the pipeline endpoint as Active
- ML studio showing the scheduled run

The screenshot shows the ML Studio Pipelines page. The 'Pipeline endpoints' tab is selected. The table below shows the 'Bankmarketing Train' pipeline endpoint, which is in the 'Active' status.

Name	Description	Date updated	Updated by	Last run submit time	Last run status	Status
Bankmarketing Train	Training bankmarketing pipeline	April 14, 2021 5:41 AM	chinmay Das	April 14, 2021 5:41 AM	Finished	Active

A large green checkmark is overlaid on the 'Active' status.

✓ ML studio showing the pipeline endpoint as Active

[DOWNLOAD PROJECT](#)

RETURN TO PATH
